

# ASSIGNMET 4

## 1. Download the dataset.

### Importing libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

## 2. loading the dataset

```
from google.colab import files
upload = files.upload()
```

```
a = pd.read_csv('abalone.csv')
```

```
a.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7

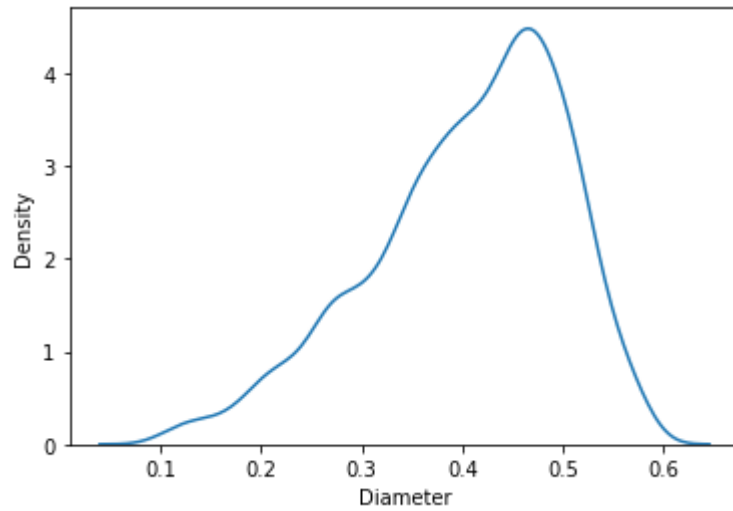
```
a['age'] = a['Rings']+1.5
a = a.drop('Rings',axis = 1)
```

## 3. Performing Visualizations.

### 1. univariate Analysis.

```
sns.kdeplot(a['Diameter'])
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36f75625d0>
```



## 2. Bi-Variate Analysis

```
ss.boxplot(x=a.Diameter,y=a.Length,palette='rainbow')
```

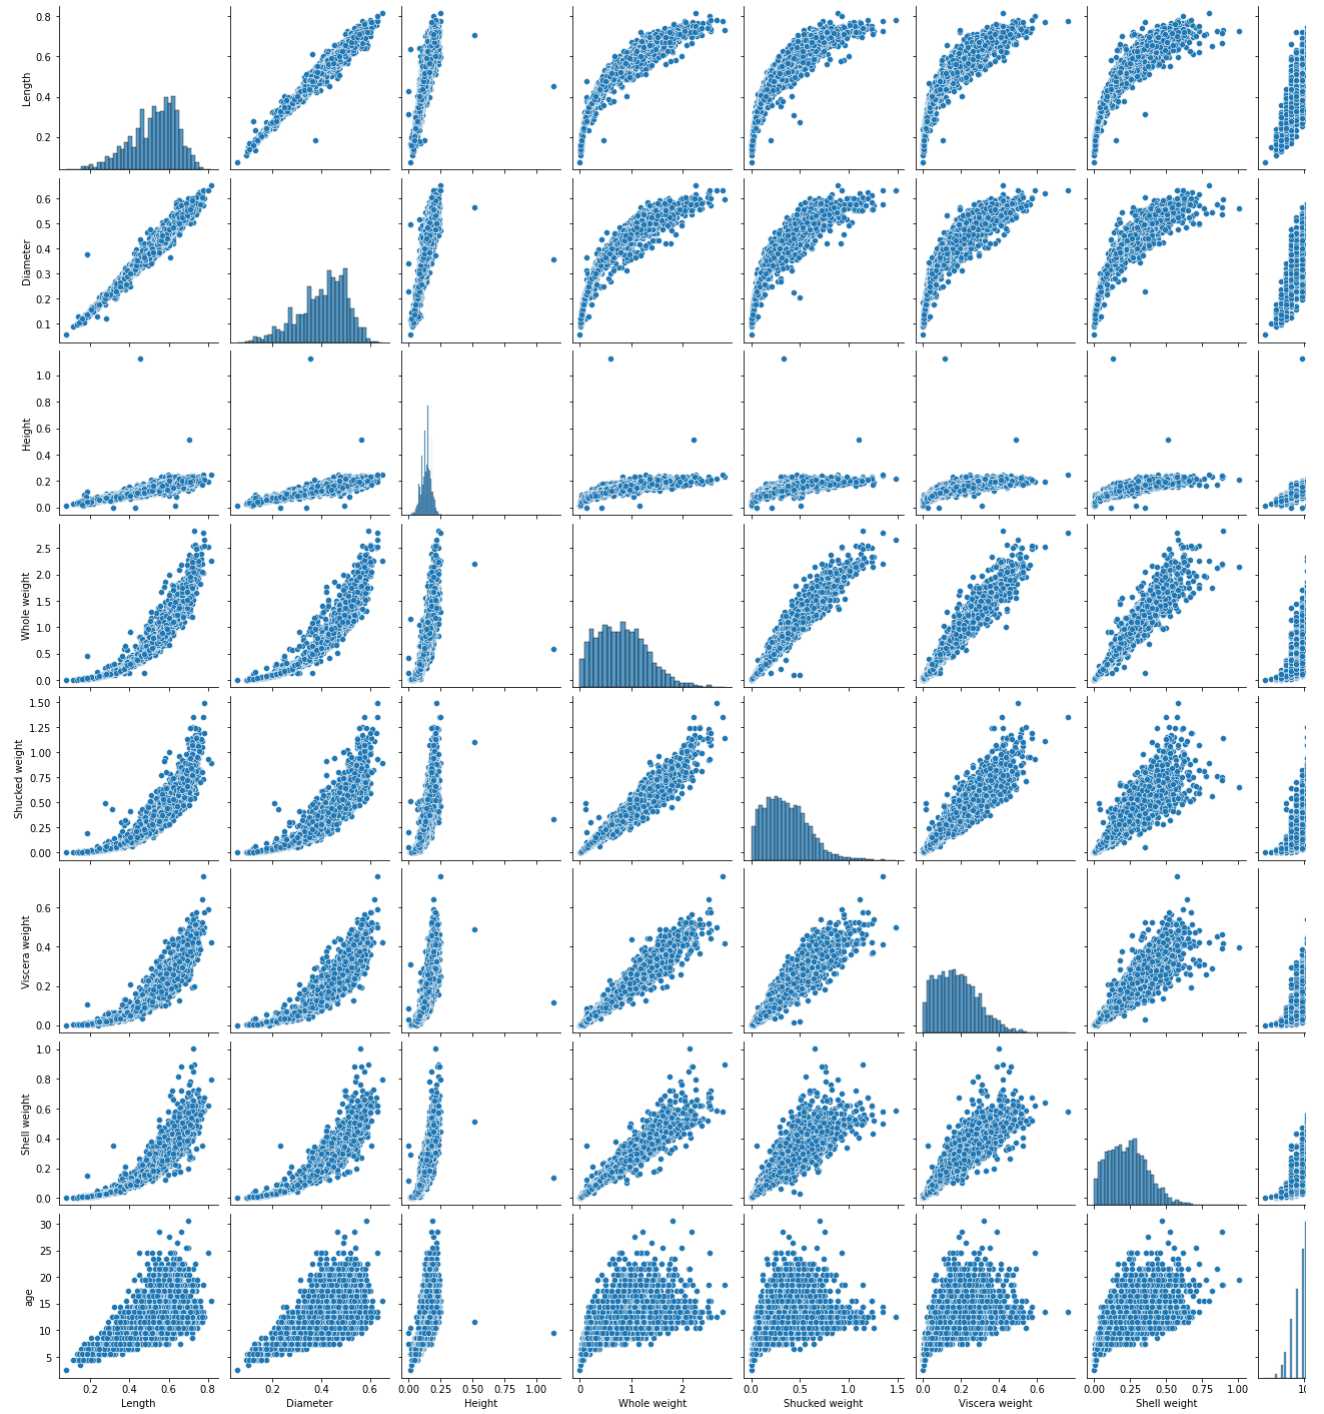
```
<matplotlib.axes._subplots.AxesSubplot at 0x7f36ff791dd0>
```



## 3. Multi-Variate Analysis

```
ss.pairplot(a)
```

<seaborn.axisgrid.PairGrid at 0x7f36ffc11550>



#### 4. Perform descriptive statistics on the dataset.

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   Sex                    4177 non-null   object  
1   Length                 4177 non-null   float64 
2   Diameter               4177 non-null   float64 
3   Height                 4177 non-null   float64 
4   Whole weight           4177 non-null   float64 
5   Shucked weight         4177 non-null   float64 
6   Viscera weight         4177 non-null   float64 
7   Shell weight           4177 non-null   float64 
8   age                    4177 non-null   float64 
dtypes: float64(8), object(1)
memory usage: 293.8+ KB
```

```
a['Diameter'].describe()
```

```
count    4177.000000
mean       0.407881
std        0.099240
min        0.055000
25%        0.350000
50%        0.425000
75%        0.480000
max        0.650000
Name: Diameter, dtype: float64
```

```
a['Sex'].value_counts()
```

```
M    1528
I    1342
F    1307
Name: Sex, dtype: int64
```

#### 5. Check for Missing values and deal with them.

```
a.isnull()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
0	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	..
4172	False	False	False	False	False	False	False	False
4173	False	False	False	False	False	False	False	False
4174	False	False	False	False	False	False	False	False
4175	False	False	False	False	False	False	False	False
4176	False	False	False	False	False	False	False	False

4177 rows x 9 columns


```
a.isnull().sum()
```

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight
Viscera weight
Shell weight 0
age          0
dtype: int64
```


## 6. Find the outliers and replace them outliers.

```
# outlier handling
a = dp.get_dummies(a)
dummy_a = a
```

```
var = 'Viscera weight'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)
```




```
a.drop(a[(a['Viscera weight'] > 0.5) &
(a['age'] < 20)].index, inplace = True)
a.drop(a[(a['Viscera weight'] < 0.5) & (
```




```
a.drop(a[(a['Shell weight'] > 0.6) &
(a['age'] < 25)].index, inplace = True)
a.drop(a[(a['Shell weight'] < 0.8) & (
a['age'] > 25)].index, inplace = True)
```

```
var = 'Shucked weight'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)
```



```
a.drop(a[(a['Whole weight'] >= 2.5) &
        (a['age'] < 25)].index, inplace = True)
a.drop(a[(a['Whole weight'] < 2.5) & (
a['age'] > 25)].index, inplace = True)
```

```
var = 'Diameter'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)
```




```
a.drop(a[(a['Diameter'] < 0.1) &
        (a['age'] < 5)].index, inplace = True)
a.drop(a[(a['Diameter'] < 0.6) & (
a['age'] > 25)].index, inplace = True)
a.drop(a[(a['Diameter'] >= 0.6) & (
a['age'] < 25)].index, inplace = True)
```

```
var = 'Height'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)
```



```
a.drop(a[(a['Height'] > 0.4) &
        (a['age'] < 15)].index, inplace = True)
a.drop(a[(a['Height'] < 0.4) & (
a['age'] > 25)].index, inplace = True)
```

```
var = 'Length'
tlf.scatter(x = a[var], y = a['age'])
tlf.grid(True)
```



```
a.drop(a[(a['Length'] < 0.1) &
        (a['age'] < 5)].index, inplace = True)
a.drop(a[(a['Length'] < 0.8) & (
a['age'] > 25)].index, inplace = True)
a.drop(a[(a['Length'] >= 0.8) & (a['age'] < 25)].index, inplace = True)
```

## 7. Check for Categorical columns and perform encoding.

```
numerical_features = a.select_dtypes(include = [pn.number]).columns
categorical_features = a.select_dtypes(include = [pn.object]).columns
```

```
numerical_features
```

```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
       'Viscera weight', 'Shell weight', 'age', 'Sex_F', 'Sex_I', 'Sex_M'],
      dtype='object')
```

```
categorical_features
```

```
Index([], dtype='object')
```

Encoding




```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
print(a.Length.value_counts())
```

```
0.550    93
0.575    93
0.625    93
0.580    92
0.600    86
..
0.755     2
0.220     2
0.150     1
0.135     1
0.760     1
Name: Length, Length: 126, dtype: int64
```

## 8. Splitting the data into dependent and independent variables.

```
x=a.iloc[:, :5]
x
```

	Length	Diameter	Height	Whole weight	Shucked weight	
<b>0</b>	0.455	0.365	0.095	0.5140	0.2245	
<b>1</b>	0.350	0.265	0.090	0.2255	0.0995	
<b>2</b>	0.530	0.420	0.135	0.6770	0.2565	
<b>3</b>	0.440	0.365	0.125	0.5160	0.2155	
<b>4</b>	0.330	0.255	0.080	0.2050	0.0895	
...	...	...	...	...	...	
<b>4172</b>	0.565	0.450	0.165	0.8870	0.3700	
<b>4173</b>	0.590	0.440	0.135	0.9660	0.4390	
<b>4174</b>	0.600	0.475	0.205	1.1760	0.5255	
<b>4175</b>	0.625	0.485	0.150	1.0945	0.5310	
<b>4176</b>	0.710	0.555	0.195	1.9485	0.9455	

4096 rows x 5 columns

```
y=a.iloc[:, :5]
y
```

	Length	Diameter	Height	Whole weight	Shucked weight
0	0.455	0.365	0.095	0.5140	0.2245
1	0.350	0.265	0.090	0.2255	0.0995
2	0.530	0.420	0.135	0.6770	0.2565
3	0.440	0.365	0.125	0.5160	0.2155
4	0.330	0.255	0.080	0.2050	0.0895
...	...	...	...	...	...
4172	0.565	0.450	0.165	0.8870	0.3700
4173	0.590	0.440	0.135	0.9660	0.4390
4174	0.600	0.475	0.205	1.1760	0.5255
4175	0.605	0.485	0.150	1.0045	0.5010



## 9. Scale the independent variables.

```
from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
x_train=ss.fit_transform(x_train)
```

```
mlrpred=mlr.predict(x_test[0:9])
```

mlrpred

```
array([[0.41  , 0.31  , 0.125 , 0.3595, 0.1415],
       [0.585 , 0.435 , 0.14  , 0.6955, 0.3085],
       [0.575 , 0.43  , 0.13  , 0.7425, 0.2895],
       [0.67  , 0.525 , 0.165 , 1.6085, 0.682 ],
       [0.645 , 0.51  , 0.2   , 1.5675, 0.621 ],
       [0.7   , 0.535 , 0.16  , 1.7255, 0.63  ],
       [0.41  , 0.325 , 0.1   , 0.3245, 0.132 ],
       [0.58  , 0.425 , 0.15  , 0.844  , 0.3645],
       [0.465 , 0.375 , 0.135 , 0.6    , 0.2225]])
```

Double-click (or enter) to edit

## 10. Splitting the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

## 11. Building the model.


```
from sklearn.linear_model import LinearRegression
mlr=LinearRegression()
mlr.fit(x_train,y_train)
```

```
LinearRegression()
```


## 12. Training the model

## 13. Testing the model

```
x_test[0:5]
```

	Length	Diameter	Height	Whole weight	Shucked weight	
<b>3268</b>	0.410	0.310	0.125	0.3595	0.1415	
<b>2668</b>	0.585	0.435	0.140	0.6955	0.3085	
<b>3042</b>	0.575	0.430	0.130	0.7425	0.2895	
<b>1040</b>	0.670	0.525	0.165	1.6085	0.6820	
<b>184</b>	0.645	0.510	0.200	1.5675	0.6210	

```
y_test[0:5]
```

	Length	Diameter	Height	Whole weight	Shucked weight	
<b>3268</b>	0.410	0.310	0.125	0.3595	0.1415	
<b>2668</b>	0.585	0.435	0.140	0.6955	0.3085	
<b>3042</b>	0.575	0.430	0.130	0.7425	0.2895	
<b>1040</b>	0.670	0.525	0.165	1.6085	0.6820	
<b>184</b>	0.645	0.510	0.200	1.5675	0.6210	

## 14. Measure the performance using Metrics.

```
from sklearn.metrics import r2_score
r2_score(mlr.predict(x_test),y_test)
```

```
1.0
```